



Northeastern


EECE 5698 Project: Movie Recommender System using Collaborative Filtering

Emmanuel Ojuba
Iuliia Klykova


Types of recommender systems

- Demographic Filtering
- Content Based filtering
- Collaborative Filtering
- Hybrid Filtering

COLLABORATIVE FILTERING



CONTENT-BASED FILTERING




Dataset

The Movies Dataset from Kaggle.com:

- The dataset contains: movies_metadata.csv, keywords.csv, credits.csv, links.csv, links_small.csv, ratings.csv, ratings_small.csv:
- **Ratings.csv** file that consists of 26 million ratings from 270,000 users for all 45,000 movies. Ratings are on a scale of 1-5.
- We picked 450,000 random samples and randomly divided them into 5 folds.
 - Fold 0: 121650 users and 8611 items.
 - Fold 1: 121573 users and 8613 items.
 - Fold 3: 121646 users and 8578 items.
 - Fold 3: 121695 users and 8610 items.
 - Fold 4: 121693 users and 8593 items



Collaborative Filtering (Matrix Factorization)



r_{ij} : rating by user i to item j .

$$r_{ij} = \langle u_i, v_j \rangle + \varepsilon_{ij} \quad , \text{ where } u_i \in \mathbb{R}^d, v_j \in \mathbb{R}^d$$

user profile

item profile

Assumption: Bilinear relationship between user profiles, u_i , and item profiles v_j

Goal: To learn u_i and v_j , given a sparse set of ratings

Objective Function: Least-Square Error with Parameter Regularization

Optimization: Stochastic Gradient Descent;
Alternating Least Squares

$$\text{RSE}(U, V) = \sum_{(i,j,r_{ij}) \in \mathcal{D}} (u_i^T v_j - r_{ij})^2 + \lambda \sum_{i=1}^n \|u_i\|_2^2 + \mu \sum_{j=1}^m \|v_j\|_2^2,$$

Alternating Least Squares

$$\text{RSE}(U, V) = \sum_{(i,j,r_{ij}) \in \mathcal{D}} (u_i^\top v_j - r_{ij})^2 + \lambda \sum_{i=1}^n \|u_i\|_2^2 + \mu \sum_{j=1}^n \|v_j\|_2^2,$$

We alternate between fixing $U \ni \{u_1, u_2, \dots, u_n\}$ and $V \ni \{v_1, v_2, \dots, v_n\}$

With either U or V fixed, the objective becomes convex least-squares estimation problem that can be optimized using the an analytic formulation:

$$v_j = \left[\sum_i u_i u_i^\top + \mu I \right]^{-1} \left[\sum_i r_{ij} u_i \right]$$

This process is allowed to continue for a set number of iterations.

We utilize Spark's MLLib's implementation which offers these hyperparameters:

- rank - Iterations - numBlocks - nonNegative - lambda

Stochastic Gradient Descent

1. Randomly initialize the U and V matrices
2. Predict the ratings for the initialized U and V matrices
3. Update U and V iteratively until convergence using a stochastic estimate of the gradient, obtained by calculating the gradient over a subsampled set of the data:



$$\nabla_{u_i} \widetilde{RSE} = 2 \sum_{v_j, \text{subsampled}} \delta_{ij} v_j + 2\lambda u_i$$

$$\nabla_{v_j} \widetilde{RSE} = 2 \sum_{u_i, \text{subsampled}} \delta_{ij} u_i + 2\mu v_j$$

$$u^{k+1} = u^k - \gamma \nabla_u \widetilde{RSE}(U, V)$$

$$v^{k+1} = v^k - \gamma \nabla_v \widetilde{RSE}(U, V)$$

Results: Alternating Least Squares (ALS)





Number of latent features $d = 2$

Regularization parameters $\lambda = \mu = 0.3$

The smallest $TestRMSE = 1.1248$

Results: Stochastic Gradient Descent





Number of latent features $d = 3$

Regularization parameters $\lambda = \mu = 10.0$

The smallest $TestRMSE = 3.388762$

Parallelism



ALS

- almost 2x speed up between 1 and 8 numbers of blocks used to parallelize computation

SGD

- no speed up observed

Conclusions

- Best performing model is the ALS-optimized MF with parameters of $d = 2$, $\lambda = \mu = 0.3$. This yielded an RMSE of 1.125
- Unexpectedly, there is huge discrepancy between the best performing ALS and SGD models, with SGD yielding an RMSE of 3.389
- Parallelism yields a speedup in the ALS computation, although not in our implementation of SGD.

Future Work

- Hybrid Approaches
- Deep Learning Methods

References

- [1] - Ahmed I. Getting Started with a Movie Recommendation System.
<https://www.kaggle.com/ibtesama/getting-started-with-a-movie-recommendation-system/notebook>
- [2] - The Movies Dataset <https://www.kaggle.com/rounakbanik/the-movies-dataset>
- [3] - Koren Y., Bell R., Volinsky C. Matrix Factorization Techniques for Recommender Systems
- [4] - Ioannidis, E. Homework 4 Handout
- [5] - Matrix Completion via Alternating Least Squares -
<http://stanford.edu/~rezab/classes/cme323/S15/notes/lec14.pdf>
- [6] - Various Implementations of Collaborative Filtering -
<https://towardsdatascience.com/various-implementations-of-collaborative-filtering-100385c6dfe0>

